

## The Beer Game — Precise Rules & Implementable Math (Python-ready)

Below is a complete, code-oriented spec for a classic four-echelon Beer Game: **Retailer → Wholesaler → Distributor → Factory**. It's written so you can hand it to Codex and implement directly. All symbols and arrays are consistent across roles and time.

---

### 1) Core Entities & Indices

- Discrete time: weeks  $t=0,1,2,\dots,T-1$
  - Roles  $r \in \{\text{Retailer}, \text{Wholesaler}, \text{Distributor}, \text{Factory}\}$
  - Each role operates the same state machine; only the **external demand source** differs:
    - Retailer's **customer demand**  $D_t$  arrives exogenously.
    - Wholesaler's demand is **Retailer's order**; Distributor's demand is **Wholesaler's order**; Factory's demand is **Distributor's order**.
- 

### 2) Parameters (per role unless noted)

- $L_{\text{ship}}$  — **shipment lead time** (integer weeks, e.g., 2). Shipments take  $L_{\text{ship}}$  periods to arrive.
- $L_{\text{prod}}$  — **production lead time** for Factory only (integer weeks, e.g., 2). Often modeled as an extra pipeline like shipping.
- $h$  — **holding cost** per unit per week (non-negative float).
- $p$  — **backlog/penalty cost** per unit per week (non-negative float).
- $I_0$  — initial on-hand inventory (non-negative int).
- $B_0$  — initial backlog (unfulfilled demand; usually 0).
- $Q_0$  — initial **incoming shipment pipeline** (length  $L_{\text{ship}}$ ); array of future arrivals each week.
- $P_0$  — initial **production pipeline** (Factory only; length  $L_{\text{prod}}$ ).
- Policy parameters (optional) for an order policy (e.g., base-stock  $SS$ , safety stock, etc.).

- Demand series  $\{D_t\}_{t=0}^{T-1}$  for Retailer (random or scripted). Upstream roles **do not** see  $D_t$  directly—only the orders they receive from their downstream partner.
- 

### 3) State Variables (per role, each week $tt$ )

- $ItI_t$  — on-hand inventory **at start** of step  $tt$ .
  - $BtB_t$  — backlog **at start** of step  $tt$  (unmet demand/orders from downstream).
  - $AtA_t$  — arrival from upstream **at start** of step  $tt$  (pops from shipment pipeline).
  - $XtX_t$  — amount **available to ship** to downstream during step  $tt$  after arrivals and before shipping (see sequence below).
  - $StS_t$  — **shipment sent** to downstream at step  $tt$  (bounded by available stock and downstream total demand).
  - $OtO_t$  — **order placed** to upstream at step  $tt$ .
  - Pipeline arrays:
    - **Ship pipeline**  $SHIPPIPE_t$  — length  $L_{ship}$ ;  $StS_t$  is pushed to the back, and the front element is  $At+1A_{t+1}$ .
    - For Factory with production lead  $L_{prod}$ : **Prod pipeline**  $PRODPIPE_t$ ; production  $MtM_t$  is pushed in, and the front element becomes **available finished goods**  $Ft+1F_{t+1}$  next step.
  - Costs:
    - **Holding cost:**  $C_{thold} = h \cdot \max(I_t, 0)$
    - **Backlog cost:**  $C_{tback} = p \cdot B_t$
    - Total role cost  $C_t = C_{thold} + C_{tback}$ . (You may add order or production costs if desired.)
- 

### 4) Weekly Step Sequence (exact order of operations)

For **each role**  $rr$ , in **each week**  $tt$ :

#### 1. Arrivals land

- Pop the front of SHIPPIPE<sub>t</sub> → A<sub>t</sub> units arrive from upstream.
- Update on-hand:  $I_t \leftarrow I_t + A_t$
- For Factory only: pop the front of PRODPIPE<sub>t</sub> → finished goods F<sub>t</sub> also land;  $I_t \leftarrow I_t + F_t$

## 2. Observe incoming demand from downstream

- Retailer: demand is D<sub>t</sub> (exogenous).
- Others: demand is **the downstream role's order** from the same week tt:  
 $d_t(r) = O_t(\text{downstream}(r))$
- Compute **total demand to fulfill** this step:  
 $TOTDEM_t = d_t(r) + B_t$  (new demand plus backlog).

## 3. Ship to downstream (satisfy demand FIFO)

- Units available to ship:  $X_t = I_t$
- Shipment sent:  $S_t = \min(X_t, TOTDEM_t)$
- Update inventory immediately:  $I_t' = I_t - S_t$
- Update backlog:  $B_{t+1} = TOTDEM_t - S_t$  (All unmet becomes backlog.)
- Push S<sub>t</sub> into the back of **downstream's SHIPPIPE**; it will arrive after L<sub>ship</sub> weeks.

## 4. Place order upstream

- Decide O<sub>t</sub> according to the chosen policy (see §7):
  - The **order is a request**, not guaranteed to be fulfilled immediately.
- Push O<sub>t</sub> into **upstream's incoming demand stream** (that role will see it as its d<sub>t</sub>).

## 5. Factory production decision (Factory only)

- Factory chooses **production start** M<sub>t</sub>. (Often  $M_t = O_t$  or policy-driven.)

- Push  $M_t M_t$  into **PRODPIPE**; finished goods become available after  $L_{\text{prod}}$  weeks.

## 6. Costs are assessed at end of week

- Set  $I_{t+1} = \max(I_t - S_t, 0)$ . (No negative inventory; it's already non-negative because we cap shipments.)
- Compute costs for the role this week:  
 $C_{\text{hold}} = h \cdot I_{t+1}$   
 $C_{\text{back}} = p \cdot B_{t+1}$   
 $C_t = C_{\text{hold}} + C_{\text{back}}$

## Notes

- **No partial backorders by choice:** backlog always carries forward; there is no lost sales in the classic game.
- **Conservation:** You can never ship more than you have on hand; all unmet demand becomes backlog.
- **Lead times** are implemented purely via FIFO queues (pipelines).

## 5) Initialization

At  $t=0$  for each role:

- $I_0 = I_0$  = given (e.g., 12).
- $B_0 = 0$ .
- $\text{SHIPPIPE}_0$  = array of length  $L_{\text{ship}}$ , each element often initialized to a steady-state flow (e.g., 4 units each), or zeros.
- Factory only:  $\text{PRODPIPE}_0$  = array of length  $L_{\text{prod}}$  (zeros or steady-state).
- Retailer: demand stream  $D_0, \dots, D_{T-1}$  given or generated.
- Upstream/downstream links wired: retailer's upstream is wholesaler; wholesaler's upstream is distributor; distributor's upstream is factory; factory's upstream is **none**.

## 6) Cost Calculations (explicit)

At the **end** of week  $t$ :

- Holding cost:

$$C_{\text{hold}} = h \cdot I_{t+1}. C^{\{\text{hold}\}}_t = h \cdot I_{t+1}.$$

- Backlog (shortage) cost:

$$C_{\text{back}} = p \cdot B_{t+1}. C^{\{\text{back}\}}_t = p \cdot B_{t+1}.$$

- Total role cost this week:

$$C_t = C_{\text{hold}} + C_{\text{back}}. C_t = C^{\{\text{hold}\}}_t + C^{\{\text{back}\}}_t.$$

- Horizon totals:

$$C_{\text{role}} = \sum_{t=0}^{T-1} C_t, C_{\text{system}} = \sum_{\text{roles}} C_{\text{role}}. C^{\{\text{role}\}} = \sum_{t=0}^{T-1} C_t, \quad C^{\{\text{system}\}} = \sum_{\{\text{roles}\}} C^{\{\text{role}\}}.$$

Optional extras you may include:

- **Order cost**  $k \cdot 1_{[O_t > 0]} \cdot \mathbb{1}_{[O_t > 0]}$  or linear per-unit order cost.
- **Capacity limits** on shipping  $S_t \leq S_{\max}$  and/or production  $M_t \leq M_{\max}$ .
- **Lost sales:** replace backlog with lost demand; change cost accordingly (not standard).

## 7) Order/Production Policies (plug-and-play)

You can implement any policy. Two common ones:

### A) Base-Stock (Order-up-to) Policy

- Define **inventory position** at time  $t$ :

$$IP_t = I_t' - B_{t+1} + \sum_{\text{on the way to me}} \text{SHIPPIE incoming}. \quad \text{IP}_t = I_t' - B_{t+1} + \underbrace{\sum_{\{\text{on the way to me}\}} \{\text{SHIPPIE incoming}\}}.$$

(Use  $I_t'$  **after shipping** this step.)

- Target level  $SS$  (per role).
- Order:

$$O_t = \max(0, SS - IP_t). O_t = \max(0, SS - IP_t).$$

- Factory production typically set to  $M_t = O_t$  (or clip to capacity).

## B) Proportional-Integral (PI) Heuristic

- Let  $e_t = \text{target\_inv} - I_{t+1} + B_{t+1}$
- $O_t = \alpha \cdot \text{dem\_forecast}_t + \beta \cdot e_t + \gamma \cdot \sum_{k=0}^t e_k$
- Clip to non-negative; integers if desired.

For the classic classroom game, base-stock is simplest and deterministic.

## 8) Exact Python-Friendly Step Function

Below is a precise, language-agnostic signature with all inputs/outputs you need for each role per week. (Codex can convert to dataclasses easily.)

### Inputs at week $t$ (for a role):

- $I$  (int),  $B$  (int)
- `ship_pipe` (deque/queue len =  $L_{\text{ship}}$ ), where `ship_pipe[0]` arrives **now**
- For Factory also: `prod_pipe` (len =  $L_{\text{prod}}$ )
- `downstream_demand` (int)  $\leftarrow D_t$  for Retailer, or the downstream role's  $O_t$
- Parameters:  $h$ ,  $p$ ,  $L_{\text{ship}}$ , ( $L_{\text{prod}}$  for Factory), any policy params/state (e.g.,  $S$ )
- Optionally: `capacity_ship`, `capacity_prod`

### Outputs:

- $I_{\text{next}}$ ,  $B_{\text{next}}$
- $S_t$  (shipment sent),  $O_t$  (order placed),  $A_t$  (arrival), ( $F_t$  for Factory)
- updated `ship_pipe_next`, (`prod_pipe_next`)
- `cost_t`

### Algorithm skeleton (non-Factory):

1.  $A_t = \text{ship\_pipe.popleft()}; I += A_t$
2.  $\text{dem} = \text{downstream\_demand}$
3.  $\text{TOTDEM} = \text{dem} + B$

4.  $S_t = \min(I, \text{TOTDEM})$  (clip by capacity if modeled)
5.  $I_{\text{prime}} = I - S_t$
6.  $B_{\text{next}} = \text{TOTDEM} - S_t$
7. **Policy:** compute  $O_t$  using  $I_{\text{prime}}$ ,  $B_{\text{next}}$ , and pipeline contents (see base-stock)
8. `ship_pipe.append( $S_t$ )` (to be received by downstream after  $L_{\text{ship}}$ )
9.  $I_{\text{next}} = I_{\text{prime}}$
10.  $\text{cost}_t = h \cdot I_{\text{next}} + p \cdot B_{\text{next}}$

#### Factory adds:

- Step 1a.  $F_t = \text{prod\_pipe.popleft}()$ ;  $I += F_t$
- Step 7a. Decide  $M_t$  (e.g.,  $M_t = O_t$ ); push `prod_pipe.append( $M_t$ )`

### 9) Worked Micro-Example (single role, single step)

Assume a non-Factory role,  $L_{\text{ship}}=2$ ,  $h=0.5$ ,  $p=1.0$ .

At start of week  $t$ :

- $I_t=8$ ,  $B_t=3$
- $\text{SHIPPIPE}_t = [4, 2]$  (so  $A_t=4$  now)
- Downstream demand this week (new) = 5
- Base-stock target  $S=15$

Step:

1. Arrive:  $A_t=4 \rightarrow I=8+4=12$
2. Demand:  $\text{dem}=5$ ,  $\text{TOTDEM}=5+3=8$
3. Ship:  $S_t = \min(12, 8) = 8$ ;  $I' = 12 - 8 = 4$ ;  $B_{t+1} = 0$
4. Inventory position:  
Incoming pipeline sum = ship\_pipe contents **after** pop but **before** appending this week's outbound to downstream (**careful: IP counts inbound to me, not outbound I send**). For this role, inbound is whatever **its upstream** has already shipped.  
Implementation detail: store **my** inbound pipeline separately; outbound pipe

belongs to the downstream role.

So  $\text{IP}_t = I - B_{t+1} + \text{sum}(\text{inbound\_to\_me})$ . Suppose inbound\_to\_me (future arrivals) now totals 22. Then  $\text{IP}_t = 4 - 0 + 2 = 6$ .  $\text{IP}_t = 4 - 0 + 2 = 6$ .

5. Order:  $O_t = \max(0, S - \text{IP}_t) = \max(0, 15 - 6) = 9$   $O_t = \max(0, S - \text{IP}_t) = \max(0, 15 - 6) = 9$
6. Costs:  $I_{t+1} = 4 \Rightarrow C_{\text{hold}} = 0.5 \cdot 4 = 2.0$   $I_{t+1} = 4 \Rightarrow C^{\text{hold}}_t = 0.5 \cdot 4 = 2.0$ ;  $B_{t+1} = 0 \Rightarrow C_{\text{back}} = 0$   $B_{t+1} = 0 \Rightarrow C^{\text{back}}_t = 0$ ; total  $C_t = 2.0$   $C_t = 2.0$ .
7. Pipe updates: push  $S_t = 8$  into **downstream's inbound** pipeline (they'll get it in 2 weeks).

Implementation tip: **Keep inbound and outbound pipelines separate** per arc to avoid double-counting. The upstream's shipment becomes your inbound in  $L_{\text{ship}}^{\text{ship}}$  weeks.

---

## 10) System Wiring (message passing per week)

At each week  $t$ :

1. **All roles** pop arrivals, update inventory, fulfill, compute  $S_t$ .
2. **Push shipments**: each role's  $S_t$  is appended to the **downstream role's inbound pipeline queue**.
3. **Compute orders**  $O_t$ , then **deliver orders** to upstream as that role's **demand input** for week  $t$ .
4. **Factory** also pushes production  $M_t$  into its prod pipeline.
5. Compute costs and log KPIs.

To avoid simultaneity issues, do it in **two passes** each week:

- **Pass A** (receive & ship): arrivals pop  $\rightarrow$  ship & compute post-ship states (but don't mutate other roles yet)
- **Pass B** (commit): push shipments to downstream pipelines; pass orders upstream; Factory pushes production.

---

## 11) KPIs & Outputs



Per role and system:

- Weekly:  $I_t$ ,  $B_t$ ,  $A_t$ ,  $S_t$ ,  $O_t$ , costs, pipeline snapshots.
  - Totals: cumulative holding, backlog, total cost; service level  $\frac{\sum_t S_t}{\sum_t \text{dem}_t}$ ; bullwhip indicators ( $\frac{\text{var}(O_t)(O_t)}{\text{var}(D_t)(D_t)}$ ).
- 

## 12) Determinism & Randomness

- If you want repeatable runs, fix a RNG seed and pre-draw the Retailer's demand series.
  - Upstream roles **must not** see the future  $D_{t+k}$ .
- 

## 13) Reference Defaults (classic classroom settings)

- $T=50$   $T = 50$  weeks
  - $L_{\text{ship}} = 2$  for all arcs
  - Factory  $L_{\text{prod}} = 2$
  - $h=0.5$ ,  $p=1.0$  (per unit-week)
  - Initial  $I_0 = 12$ ,  $B_0 = 0$
  - Inbound pipelines initialized with steady 4 units/week (optional)
  - Retailer demand  $D_t$ : 4 units/week for weeks 0–3, then 8 thereafter (the “shock”) — or use your own series.
- 

## 14) Minimal Data Structures (Python hint)

- RoleState:  $I$ ,  $B$ , inbound\_pipe: deque[int], optional prod\_pipe: deque[int] (Factory only), policy params, logs.
- Edge: downstream inbound pipeline is a deque residing with the **downstream** role.
- One step(role\_state, downstream\_demand) function per role that returns (state\_next,  $S_t$ ,  $O_t$ , cost\_t).
- Week loop orchestrator that:

1. collects arrivals for all roles,
  2. calls step (compute but **stage** shipments and orders),
  3. commits staged shipments and orders to the appropriate queues/inputs,
  4. accumulates costs and logs.
- 

## 15) Common Pitfalls to Avoid

- **Shipping more than on-hand:** always cap  $S_t$  by available inventory after arrivals this week.
  - **Using outbound pipeline in IP:** inventory position includes **inbound to me**, not what I've promised to someone else.
  - **Order timing:** orders placed at week  $tt$  become the upstream role's **demand in the same week  $tt$**  (information lead time = 0), but **physical shipments** respect  $L_{\text{ship}}$ .
  - **Factory creation of goods:** production enters  $\text{prod\_pipe}$  and only becomes inventory after  $L_{\text{prod}}$  weeks.
  - **Backlog arithmetic:** always  $B_{t+1} = \text{previous\_backlog} + \text{new\_demand} - \text{shipped}$  (never negative).
- 

If you want, I can turn this spec into a clean Python scaffold (dataclasses + deterministic simulator + pluggable policies) in a single file.